



Compiling Fundamentals

15-213/15-513/14-513: Introduction to Computer Systems



Questions that will be answered today

- What does it mean to compile code?
- What does compiling code look like?
- How can code be compiled?
- What are Makefiles?



Why is this important?

- It is important to understand how programs are compiled to have a better understanding of how different parts of a computer interact with each other.
- Fundamental aspect of how computers run code.


Levels of Abstraction

- C [and other high level languages] are easy for programmers to understand, but computers require lots of software to process them
- Machine code is just the opposite: easy for the computer to process, humans need lots of help to understand it
- Assembly language is a compromise between the two: readable by humans (barely), close correspondence to machine code


C programmer

```
#include <stdio.h>
int main(){
    int i, n = 10, t1 = 0, t2 = 1, nxt;
    for (i = 1; i <= n; ++i){
        printf("%d, ", t1);
        nxt = t1 + t2;
        t1 = t2;
        t2 = nxt; }
    return 0; }
```


Assembly programmer



Computer designer




Gates, clocks, circuit layout, ...



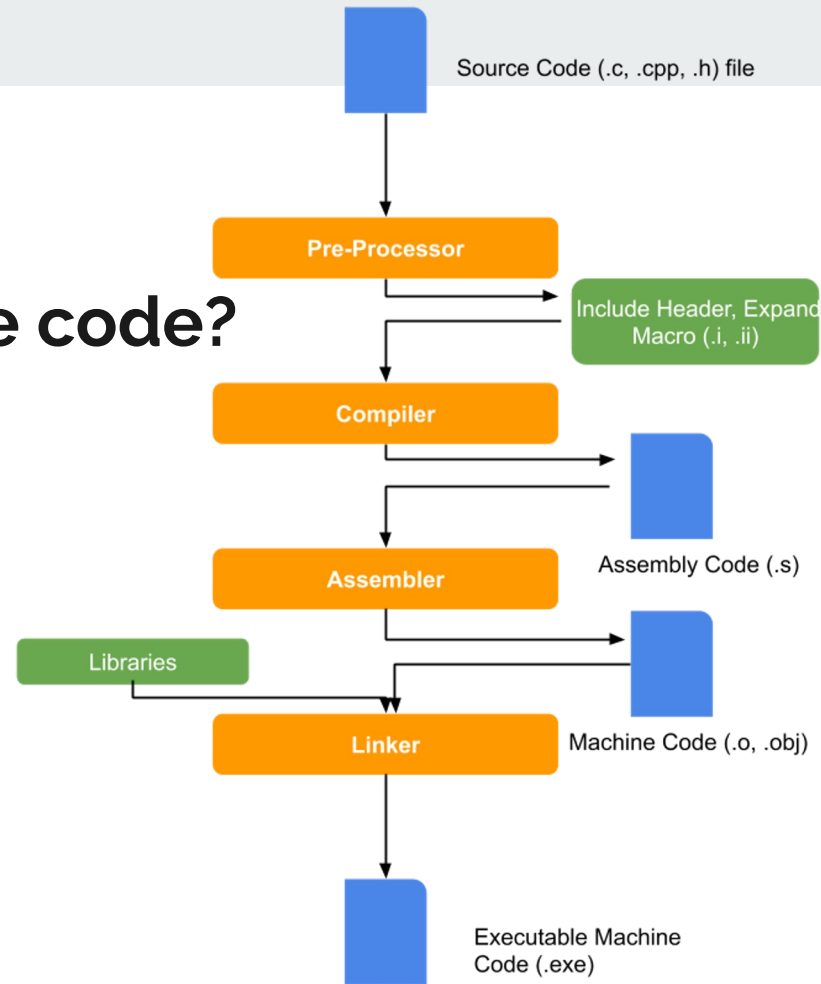
What does it mean to compile code?

- The computer only understands *machine code* directly
- All other languages must be either
 - *interpreted*: executed by software
 - *compiled*: translated to machine code by software



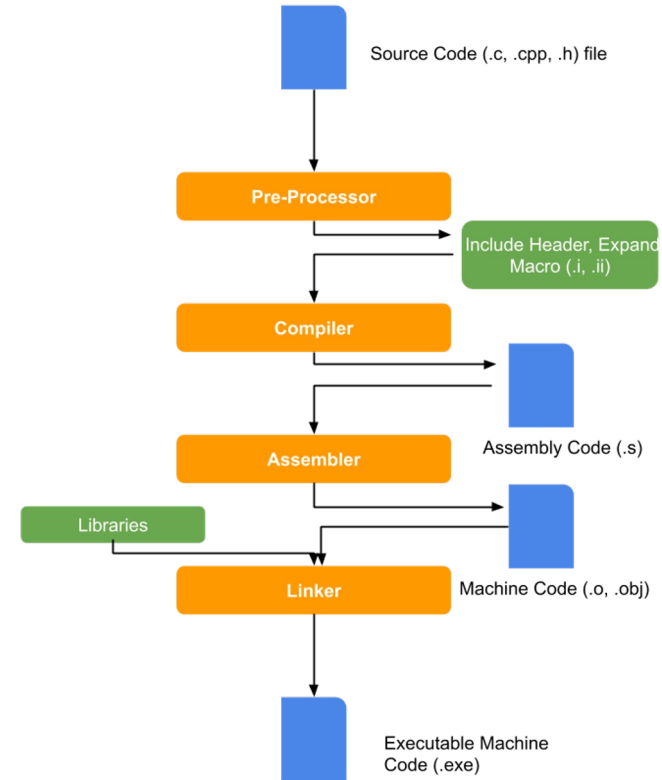
What does it mean to compile code?

- Computer follows steps to **translate** your code into something the computer can understand
- This is the process of **compiling** code [a compiler completes these actions]
- Four steps for C: preprocessing, compiling, assembling, linking
 - Most other compiled languages don't have the preprocessing step, but do have the other three



Stepping through the stages

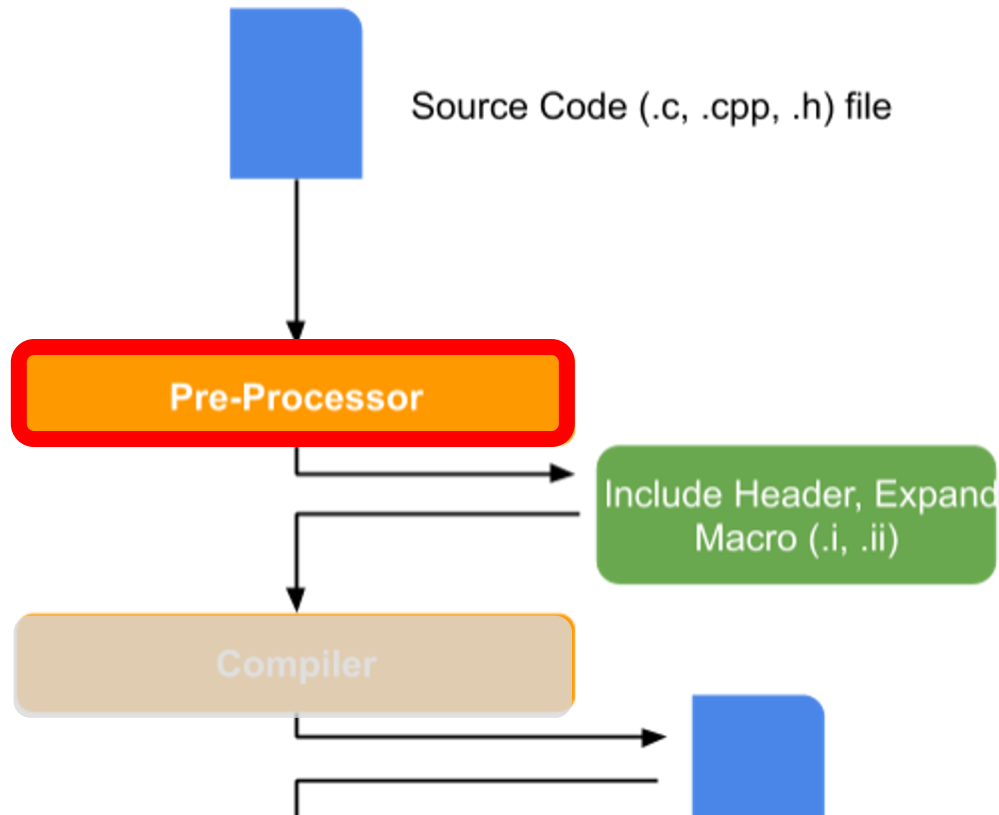
- Pre-Processor
 - `$ gcc -E [flags] [filenames]`
- Compiler
 - `$ gcc -S [flags] [filenames]`
- Assembler
 - `$ gcc -c [flags] [filenames]`
 - `$ objdump -d [filenames]`
- Linker
 - `$ gcc -o [exename] [flags] [filenames]`



C Code to Machine Code

Pre-Processor

- Unique to the C family; other languages don't have this
- Processes `#include`, `#define`, `#if`, macros
 - Combines main source file with headers (textually)
 - Defines and expands macros (token-based shorthand)
 - Conditionally removes parts of the code (e.g. specialize for Linux, Mac, ...)
- Removes all comments
- Output looks like C still



Before and after preprocessing

```
#include <limits.h>
#include <stdio.h>

int main(void) {
    // Report the range of `char` on this system
    printf("CHAR_MIN = %d\n"
           "CHAR_MAX = %d\n",
           CHAR_MIN, CHAR_MAX);
    return 0;
}
```

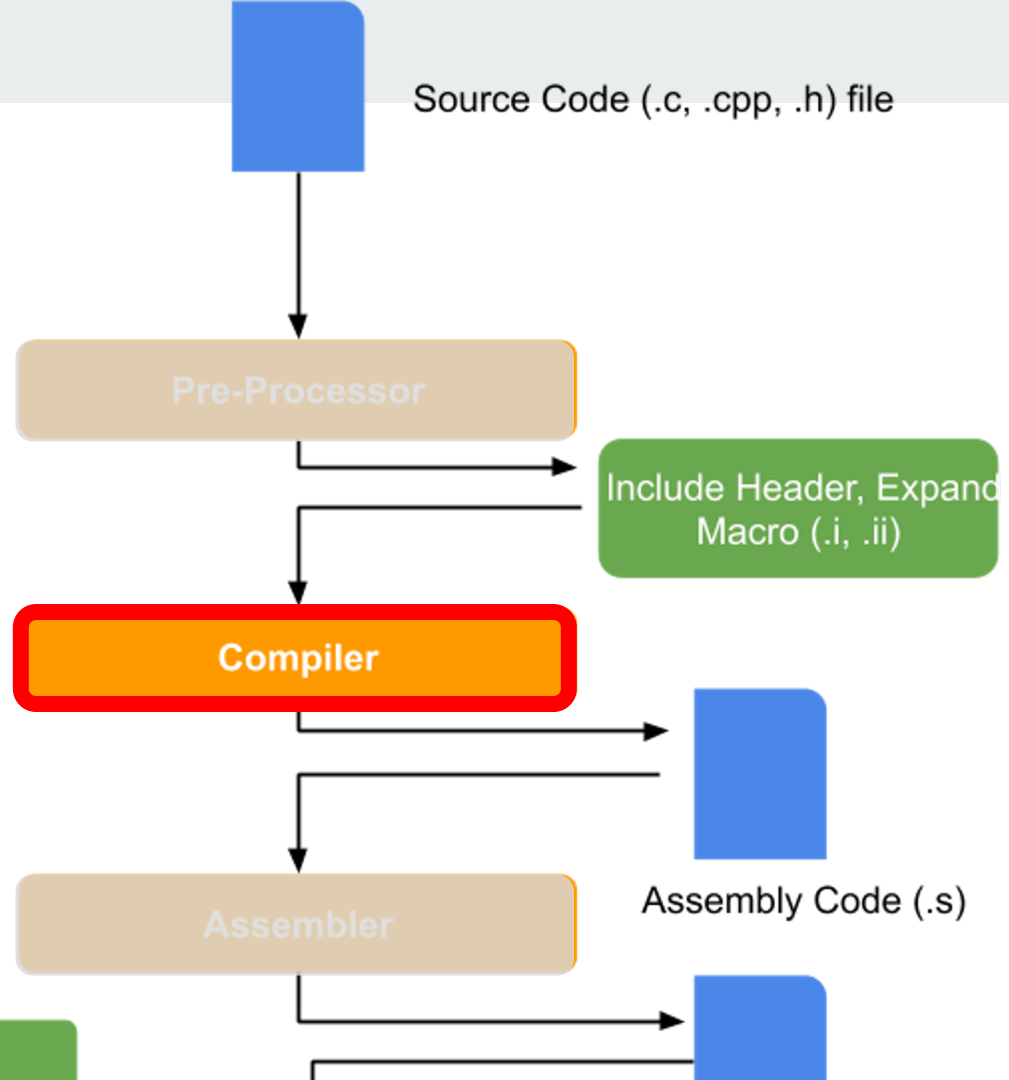
- Contents of header files inserted inline
- Comments removed
- Macros expanded
- “Directive” lines (beginning with #) communicate things like original line numbers

```
# 1 "test.c"
# 1 "/usr/lib/gcc/x86_64-linux-gnu/10/include/limits.h" 1 3 4
...
# 1 "/usr/include/stdio.h" 1 3 4
...
extern int fprintf (FILE *__restrict __stream,
                    const char *__restrict __format, ...);
extern int printf (const char *__restrict __format, ...);
...
# 874 "/usr/include/stdio.h" 3 4
# 3 "test.c" 2

int main(void) {
    printf("CHAR_MIN = %d\n"
           "CHAR_MAX = %d\n",
# 6 "test.c" 3 4
           (-0x7f - 1)
# 6 "test.c"
           , 0x7f);
    return 0;
}
```

Compiler

- The compiler translates the preprocessed code into assembly code
 - This changes the format and structure of the code but preserves the semantics (what it does)
 - Can change lots of details for optimization, as long as the overall effect is the same



Before and after compilation

```
extern int printf (const char *__restrict
                  __format, ...);

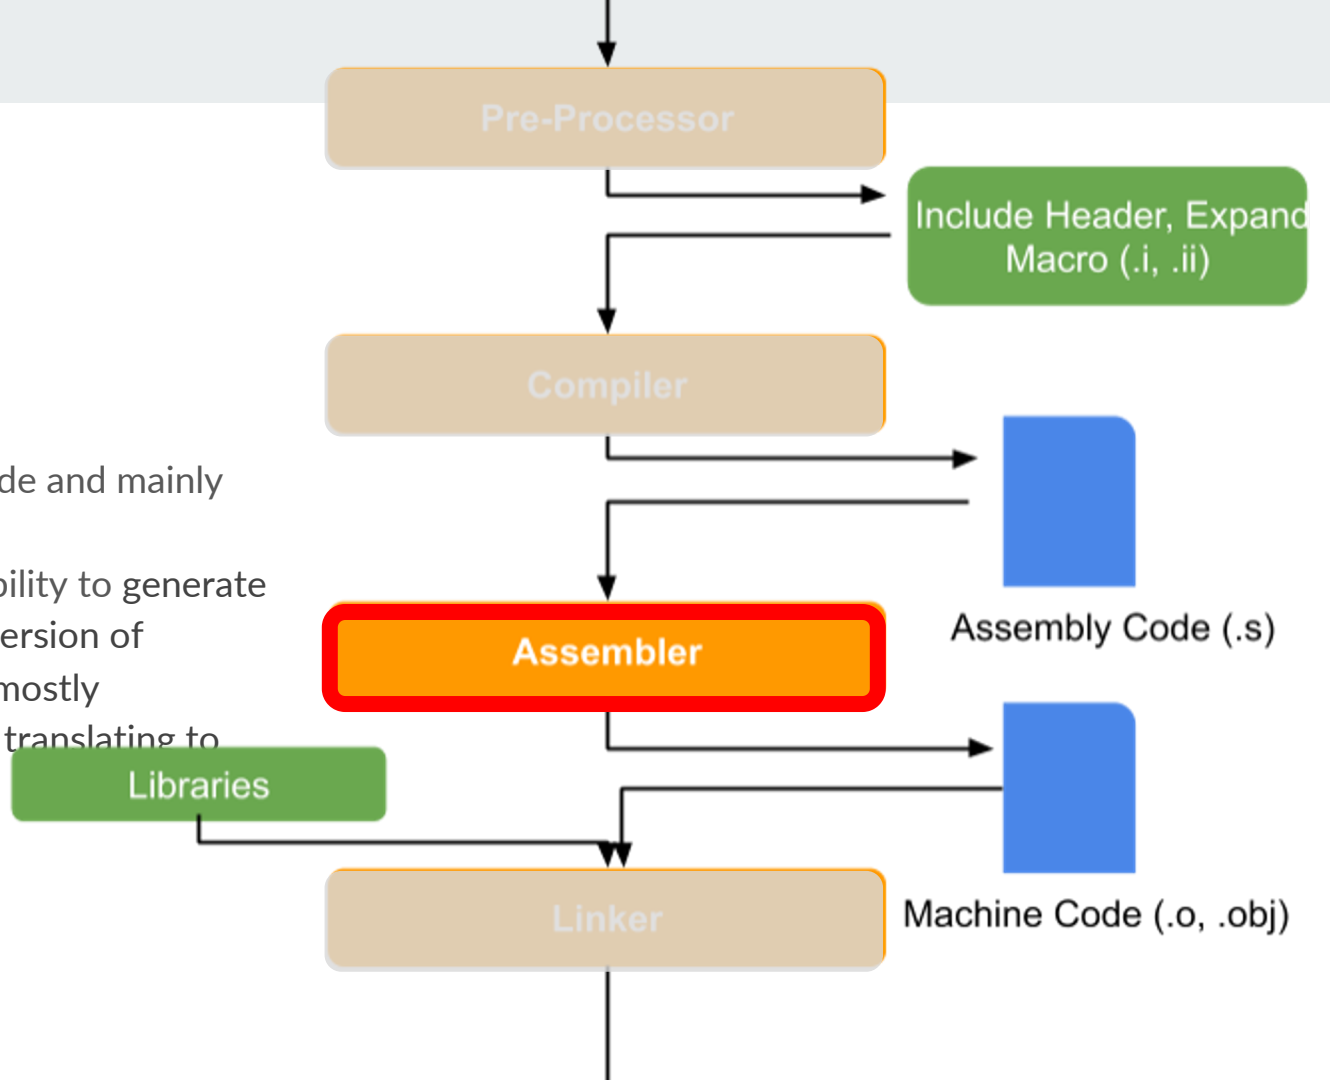
int main(void) {
    printf("CHAR_MIN = %d\n"
          "CHAR_MAX = %d\n",
          (-0x7f - 1), 0x7f);
    return 0;
}
```

- C source code converted to *assembly language*
- Textual, but 1:1 correspondence to machine language
- printf just referred to, not declared

```
.file      "test.c"
.section   .rodata.str1.1,"aMS",@progbits,1
.LC0:
.string    "CHAR_MIN = %d\nCHAR_MAX = %d\n"
.text
.globl     main
main:
    subq    $8, %rsp
    movl    $127, %edx
    movl    $-128, %esi
    leaq    .LC0(%rip), %rdi
    xorl    %eax, %eax
    call    printf@PLT
    xorl    %eax, %eax
    addq    $8, %rsp
    ret
.size      main, .-main
```

Assembler

- Parses assembly code and mainly translates into bits
- There is some flexibility to generate the most efficient version of machine code, but mostly responsible for just translating to bits.



Before and after assembling

```
.file      "test.c"
.section   .rodata.str1.1,"aMS",@progbits,1
.LC0:
.string    "CHAR_MIN = %d\nCHAR_MAX = %d\n"
.text
.globl     main
main:
    subq    $8, %rsp
    movl    $127, %edx
    movl    $-128, %esi
    leaq    .LC0(%rip), %rdi
    xorl    %eax, %eax
    call    printf@PLT
    xorl    %eax, %eax
    addq    $8, %rsp
    ret
.size     main, .-main
```

```
$ objdump -s -r test.o
test.o:          file format elf64-x86-64
```

```
RELOCATION RECORDS FOR [.text]:
OFFSET          TYPE          VALUE
0000000000000011 R_X86_64_PC32   .LC0-0x0000000000000004
0000000000000018 R_X86_64_PLT32  printf-0x0000000000000004
```

```
Contents of section .rodata.str1.1:
0000 43484152 5f4d494e 203d2025 640a4348  CHAR_MIN = %d.CH
0010 41525f4d 4158203d 2025640a 00          AR_MAX = %d..
```

```
Contents of section .text:
0000 4883ec08 ba7f0000 00be80ff ffff488d  H.....H.
0010 3d000000 0031c0e8 00000000 31c04883  =...1.....1.H.
0020 c408c3                                     ...
```

- Everything is now binary

Before and after assembling

```
.file "test.c"
.section .rodata.str1.1,"aMS",@progbits,1
.LC0:
.string "CHAR_MIN = %d\nCHAR_MAX = %d\n"
.text
.globl main
main:
    subq    $8, %rsp
    movl    $127, %edx
    movl    $-128, %esi
    leaq    .LC0(%rip), %rdi
    xorl    %eax, %eax
    call    printf@PLT
    xorl    %eax, %eax
    addq    $8, %rsp
    ret
.size     main, .-main
```

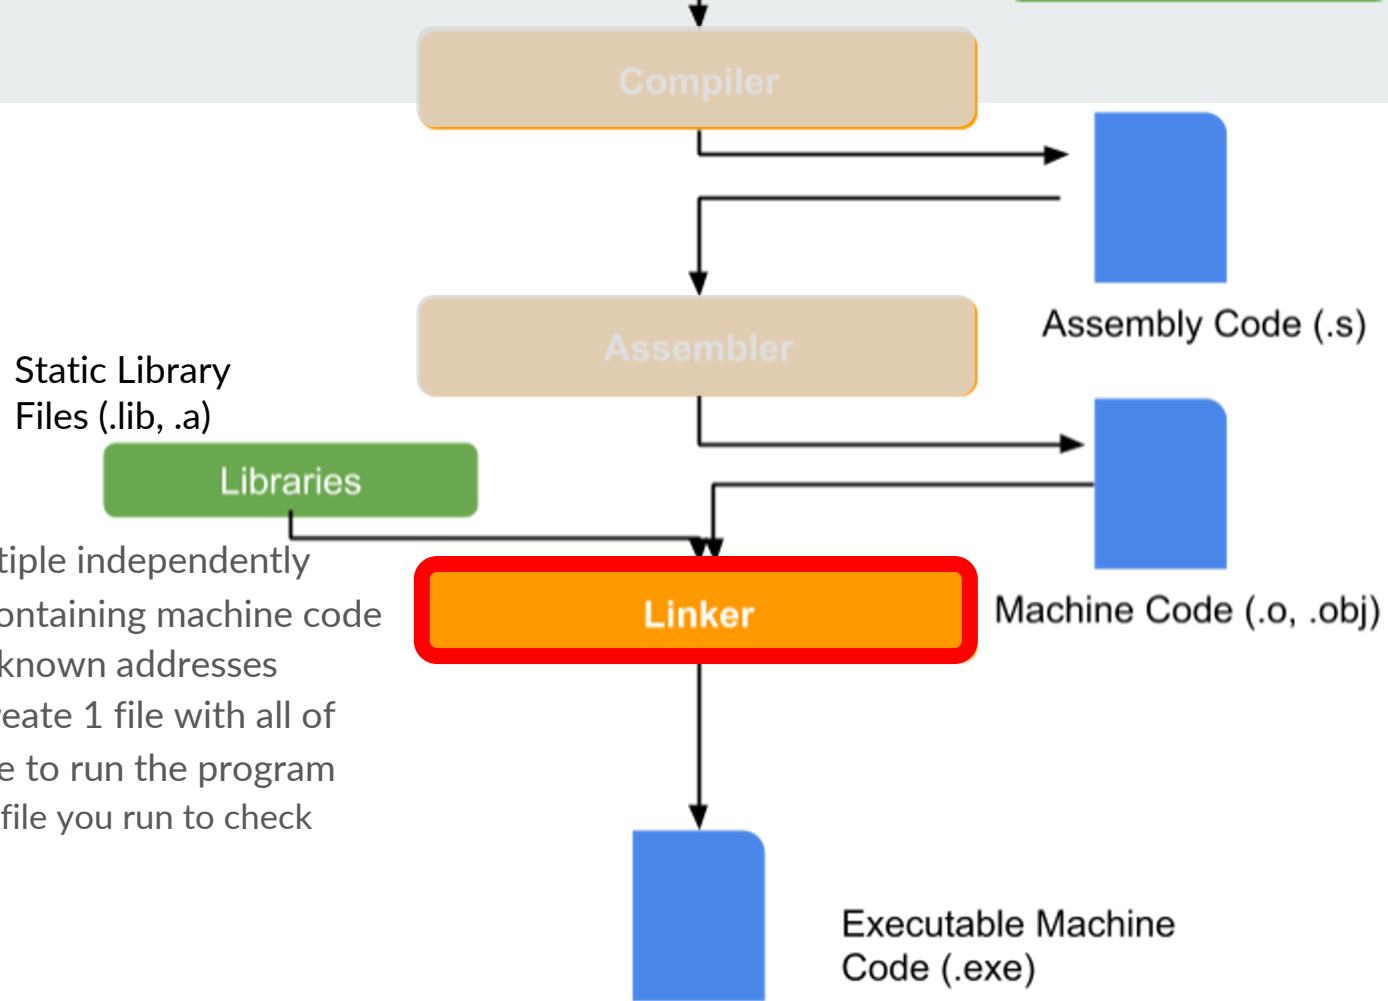
```
$ objdump -d -r test.o
test.o:          file format elf64-x86-64
Disassembly of section .text.startup:
```

```
0000000000000000 <main>:
0:  48 83 ec 08          sub     $0x8,%rsp
4:  ba 7f 00 00 00      mov     $0x7f,%edx
9:  be 80 ff ff ff      mov     $0xffffffff80,%esi
e:  48 8d 3d 00 00 00    lea     0x0(%rip),%rdi
                        11: R_X86_64_PC32 .LC0-0x4
15:  31 c0               xor     %eax,%eax
17:  e8 00 00 00 00      call    1c <main+0x1c>
                        18: R_X86_64_PLT32 printf-0x4
1c:  31 c0               xor     %eax,%eax
1e:  48 83 c4 08          add     $0x8,%rsp
22:  c3                 ret
```

- Just to emphasize that 1:1 correspondence between assembly and machine instructions

Linker

- Aggregates multiple independently compiled files containing machine code
- Fills in those unknown addresses
- The goal is to create 1 file with all of the needed code to run the program
 - This is the file you run to check your code!



How to Use The Compiler (gcc)



GCC - What is it?

- GNU Compiler Collection
 - GCC is a set of compilers for various languages. It provides all of the infrastructure for building software in those languages from source code to assembly.
- The compiler can handle compiling everything on its own, but you can use various flags to breakdown the compilation steps
- Example:

```
gcc [flags] [infile(s)]
```



Common GCC Flags

- o [EXECUTABLE NAME] : names executable file
- Ox : Code optimization
 - O0 : Compile as fast as possible, don't optimize [this is the default]
 - O1, -O2, -O3: Optimize for reduced execution time [higher numbers are more optimized]
 - Os : Optimize for code size instead of execution time.
 - Og : Optimize for execution time, but try to avoid making interactive debugging harder.
- g : produce “debug info”: annotate assembly so gdb can find variables and source code
- Wall : enable many “warning” messages that *should* be on by default, but aren't
 - Does *not* turn on all of the warning messages GCC can produce.
 - See <https://gcc.gnu.org/onlinedocs/gcc-4.8.0/gcc/Warning-Options.html> for many more
- Werror : turns all warnings into errors
- std=c99 : use the 1999 version of the C standard and disable some (not all!) extensions

Makefiles

What is a makefile?

- Automates the process of creating files (using a compiler)
- For example, create **bomb** from `bomb.c`, `phases.c`, and `util.c`
- Running `make bomb` will update `bomb`
 - Only if any of the source files have changed; avoids unnecessary work
 - Remembers complicated compiler commands for you
- Can also store recipes for automating development tasks
 - `make format` to reformat source files



Makefile



Makefiles are lists of *rules*

- There are two kinds of rules: **normal** and **phony**
 - Normal rules create files
 - Phony rules don't directly create files
- Each rule has a **target**.
 - For **normal** rules, the target is the name of the file that the rule will create
 - For **phony** rules, the target is an arbitrary name for what the rule does
- Rules may have **prerequisites** (also known as **dependencies**)
 - Prerequisites are the files that are needed to create the target
 - If any of the prerequisites doesn't exist, it must be created first
 - If any of the prerequisites is newer than the target, the target is "out of date" and must be re-created
- Rules may have **commands**.
 - One or more shell commands that create the target from its prerequisites
 - For phony rules, just some commands to be run



Normal rule example

```
bomb: bomb.o phases.o util.o
      $(CC) -o bomb bomb.o phases.o
      util.o
```

Normal rule example

If this file
doesn't exist...

... or if it is older than any of these files...

```
bomb: bomb.o phases.o util.o
      $(CC) -o bomb bomb.o phases.o
util.o
```

... then run this command.

Normal rule example

If this file
doesn't exist...

... or if it is older than any of these files...

bomb: bomb.o phases.o util.o

`$(CC) -o bomb bomb.o phases.o`

util.o

This refers to the value of a
variable, named CC, that holds
the name of a C compiler.

... then run this command.



Normal rule without prerequisites

```
output_dir:  
    mkdir output_dir
```

- Run `mkdir output_dir` if `output_dir` does not exist
- If it does exist, no action



Normal rule without commands

```
bomb.o: bomb.c support.h phases.h
```

- Re-create bomb.o if any of bomb.c, support.h, phases.h is newer
- The commands to do this are given somewhere else
 - A *pattern rule* elsewhere in the Makefile
 - An *implicit rule* built into Make



Pattern and implicit rules

`%.o: %.c`

`$(CC) $(CFLAGS) -c -o $@ $<`

- To create an .o file from a .c file with the same base name, use this command
- Special variables \$@ and \$< give the name of the .o and .c files respectively
- Variables CC and CFLAGS can be set to customize behavior
- This rule is *implicit* — built into Make — you don't have to write it yourself



Phony rule example

```
all: bomb bomb-solve  
.PHONY: all
```

- When asked to create “all”, create bomb and bomb-solve
- Does **not** create a file named “all”
- The .PHONY annotation can be anywhere in the makefile



Phony rule example 2

```
clean:
```

```
    rm -f bomb bomb-solve *.o
```

```
.PHONY: clean
```

- When asked to create “clean”, run this command
 - Which deletes bomb, bomb-solve, and all object files
- Does **not** create a file named “clean”



The `make` command

- Running `make` in the shell will cause the shell to look for a Makefile in the current directory. If it finds one, it will attempt to create the first target listed in the Makefile.
- You can also run `make <target_name>` to indicate exactly which target you want to create.
- By convention, the first target is a phony target named `all`
 - so `make` and `make all` do the same thing
 - as the name implies, this is to create everything that the makefile knows how to create
- Phony rules serve as entry points into the Makefile
 - `make all` creates everything, `make clean` deletes all generated files, `make check` runs tests, ...
 - But you can also make `bomb.o` if that's the only thing you want



A complete Makefile

```
CC = gcc
CFLAGS = -std=c99 -g -O2 -Wall -Werror

all: bomb bomb-solve
bomb: bomb.o phases.o util.o
    $(CC) $(LDFLAGS) -o $@ $^ $(LIBS)

bomb-solve: bomb.o phases-solve.o util.o
    $(CC) $(LDFLAGS) -o $@ $^ $(LIBS)

bomb.o: bomb.c phases.h support.h
phases.o: phases.c phases.h support.h
phases-solve.o: phases-solve.c phases.h support.h
util.o: util.c support.h

clean:
    rm -f bomb bomb-solve *.o

.PHONY: all clean
```

- OK to use undefined variables
 - LDFLAGS, LIBS
 - Found in environment or treated as empty
- Don't need to give commands to create object files from C source
 - But do need to list header file dependencies for each object file
- *Do* need to give commands to create executables (missing feature)
- all rule at the top, clean rule at the bottom
- One .PHONY annotation for all phony rules

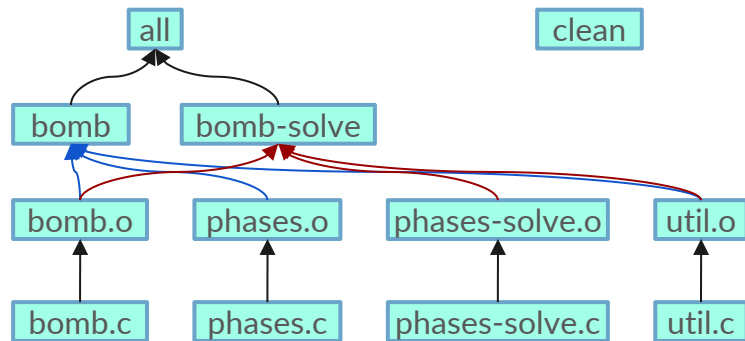
Rules form a graph

```
CC = gcc
CFLAGS = -std=c99 -g -O2 -Wall -Werror

all: bomb bomb-solve
bomb: bomb.o phases.o util.o
    $(CC) $(LDFLAGS) -o $@ $^ $(LIBS)
bomb-solve: bomb.o phases-solve.o util.o
    $(CC) $(LDFLAGS) -o $@ $^ $(LIBS)
bomb.o: bomb.c phases.h support.h
phases.o: phases.c phases.h support.h
phases-solve.o: phases-solve.c phases.h support.h
util.o: util.c support.h

clean:
    rm -f bomb bomb-solve *.o

.PHONY: all clean
```



- Make avoids unnecessary work
 - If `bomb.c` changes, make will re-create `bomb.o`, `bomb`, `bomb-solve`
 - If `phases.c` changes, make will only re-create `phases.o` and `bomb`
- Make can see through missing targets
 - If `bomb.o` does not exist, make will create it from `bomb.c`

Practice!

https://www.cs.cmu.edu/~213/bootcamps/lab3_handout.pdf

Feedback:

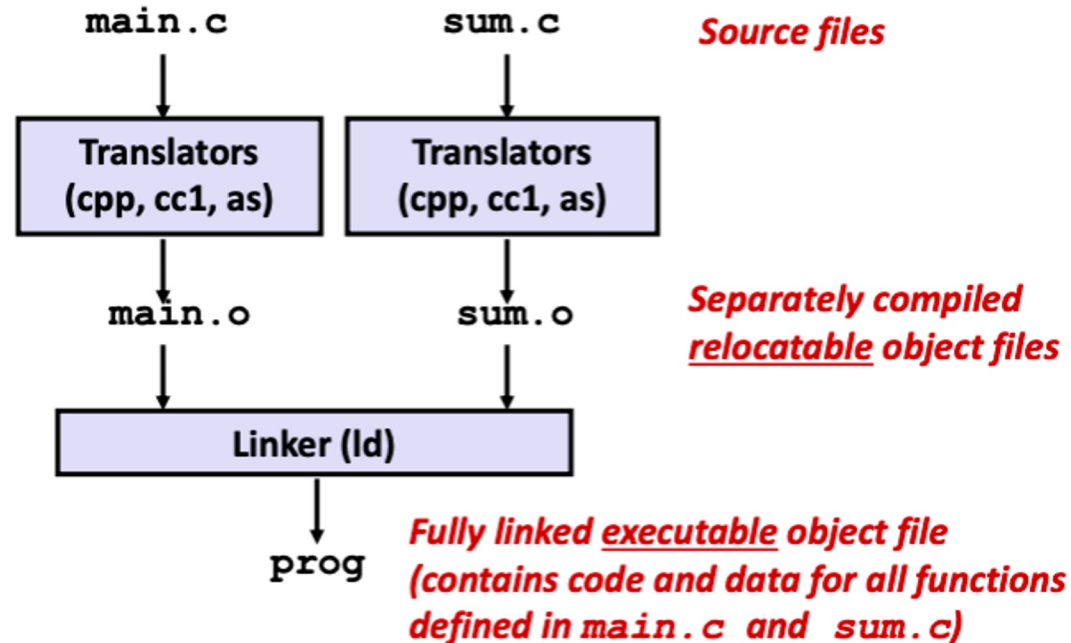
<https://tinyurl.com/213bootcamp2>

Appendix

Linking Files

Why are we learning about linking files?

- Linker is a computer system program that takes object files (generated by a compiler or an assembler) and combines them into a single executable file, library file, or another object file.
- Programs are translated and linked using a compiler driver:
 - `linux> gcc -Og -o prog main.c sum.c`
 - `linux> ./prog`
- More in future lecture!





What does a linker do?

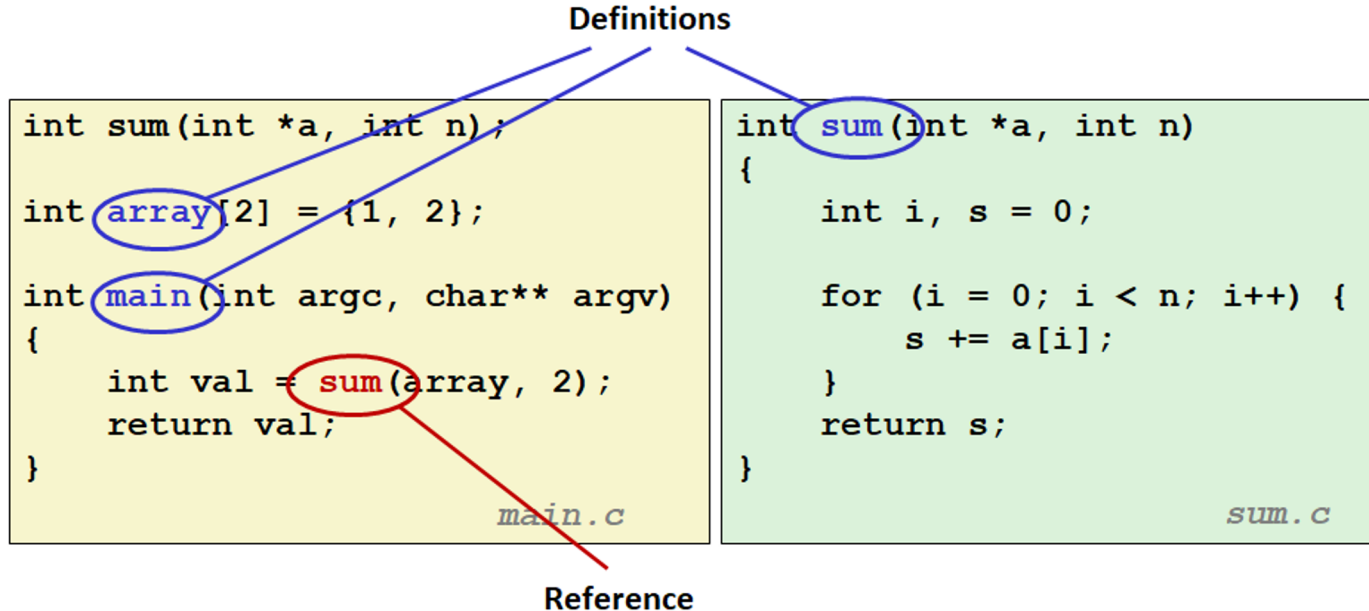
- Symbol resolution
 - Programs define and reference *symbols* (global variables and functions)
 - Linker associates each symbol reference with exactly 1 symbol definition
- Relocation
 - Merges separate code and data sections into single sections
 - Relocates symbols from relative locations in .o files to final memory locations
 - Updates all references to symbols to reflect new positions



Linker symbols

- Global symbols
 - Symbols defined by module *m* that can be referenced by other modules.
 - e.g., non-static C functions and non-static global variables.
- External symbols
 - Global symbols that are referenced by module *m* but defined by some other module.
- Local symbols
 - Symbols that are defined and referenced exclusively by module *m*.
 - e.g., C functions and global variables defined with the static attribute.
 - Local linker symbols are not local program variables

Symbols





Why do you need linkers?

- Modularity
 - Program can be written as a collection of smaller source files, rather than one monolithic mass.
- Efficiency
 - Time: Separate compilation
 - Change one source file, compile, and then relink. No need to recompile other source files.
 - Space: Libraries
 - Common functions can be aggregated into a single file...



Static vs Dynamic Linking

- Static Linking
 - Executable files and running memory images contain only the library code they actually use
- Dynamic linking
 - Executable files contain no library code
 - During execution, single copy of library code can be shared across all executing processes



Types of object files

- Relocatable object file (.o file)
 - Code and data that can be combined with other relocatable object files to form executable object file
 - Each .o file is produced from exactly one source (.c) file
- Executable object file (a.out file)
 - Code and data that can be copied directly into memory and then executed
- Shared object file (.so file)
 - Special type of relocatable object file that can be loaded into memory and linked dynamically, at either load time or run-time

How Linker resolves duplicate symbol definitions

- Program symbols are either *strong* or *weak*
 - **Strong**: procedures and initialized globals
 - **Weak**: uninitialized globals
 - Or one's declared with specifier `extern`

p1.c

```
strong → int foo=5;
strong → p1() {
}
```

p2.c

```
int foo; ← weak
p2() { ← strong
}
```



Symbol rules

1. Multiple strong symbols are not allowed
 - Each item can be defined only once
 - Otherwise, linker error
2. Given a strong symbol and multiple weak symbols, choose the strong symbol
 - References to the weak symbol resolve to the strong symbol
3. If there are multiple weak symbols, pick an arbitrary one



LD_LIBRARY_PATH

- If you are using dynamic libraries, you need to tell the compiler where to look for the library!
- It is easiest to use dynamic libraries with makefiles, just include this line:

```
LD_LIBRARY_PATH = "~/my/path"
```

- If you are interested in creating a dynamic library, follow the steps here:
 - Shared Libraries: <https://tldp.org/HOWTO/Program-Library-HOWTO/shared-libraries.html>
 - Dynamic Libraries: <https://tldp.org/HOWTO/Program-Library-HOWTO/dl-libraries.html>



Resources

<https://missing.csail.mit.edu/2020/metaprogramming/>

<https://www.cs.cmu.edu/~15131/f17/topics/makefiles/>

https://www.gnu.org/software/make/manual/html_node/Phony-Targets.html

<https://makefiletutorial.com/>

<https://www.oreilly.com/library/view/programming-embedded-systems/0596009836/ch04.html>

<https://gcc.gnu.org/onlinedocs/gcc/>

<https://daveparillo.github.io/cisc187-reader/build-tools/make.html>